

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
from tqdm import tqdm
```

Convex Optimization - Homework 3

Note : The present notebook can be found here : <https://github.com/RSLLS/ConvexOptHW3>

$$\text{minimize w.r.t. } w \quad \frac{1}{2} \|Xw - y\|_2^2 + \lambda \|w\|_1 \quad (\text{LASSO})$$

Question 1 : Derive the dual problem of LASSO and format it as a general Quadratic Problem

One get by reformulating (LASSO) :

$$\begin{aligned} &\text{minimize w.r.t. } w, z \quad \frac{1}{2} \|z\|_2^2 + \lambda \|w\|_1 \\ &\text{subject to} \quad z = Xw - y \end{aligned}$$

Let \mathcal{L} be the lagrangien of the latter problem :

$$\begin{aligned} \mathcal{L}(w, z, \eta) &= \frac{1}{2} \|z\|_2^2 + \lambda \|w\|_1 + {}^t\eta(Xw - y - z) \\ \mathcal{L}(w, z, \eta) &= \underbrace{\frac{1}{2} \|z\|_2^2 - {}^t\eta z}_{f(z)} + \underbrace{\lambda \|w\|_1 + {}^t\eta Xw - {}^t\eta y}_{g(w)} \end{aligned}$$

The next step is to minimize \mathcal{L} with respect to w and z . To do so, one needs to minimize both f and g defined above.

First, $\nabla f(z) = z - \eta$, so $\nabla f(z) = 0 \Rightarrow z = \eta$ and $f(\eta) = \frac{-1}{2} \|\eta\|_2^2$. Hence, $\inf_z f(z) = \frac{-1}{2} \|\eta\|_2^2$.

Second, $g(w) = \lambda \|w\|_1 + {}^t\eta Xw = -[{}^t\eta Xw - \lambda \|w\|_1]$. Hence :

$$\inf_w g(w) = -\lambda \sup_w \left\{ -\frac{1}{\lambda} {}^t\eta Xw - \|w\|_1 \right\}$$

According to the second homework, the conjugate of $\|\cdot\|_1$ is :

$$y \mapsto \sup_x \{ {}^t y x - \|x\|_1 \} = \begin{cases} 0 & \text{if } \|y\|_\infty \leq 1 \\ +\infty & \text{otherwise} \end{cases}$$

The latter function evaluted with $y = -\frac{1}{\lambda} {}^t X \eta$ gives exactly $\inf_w g(w)$:

$$\inf_w g(w) = \begin{cases} 0 & \text{if } \|{}^t X \eta\|_\infty \leq \lambda \\ -\infty & \text{otherwise} \end{cases}$$

All in all :

$$\inf_{w, z} \mathcal{L}(w, z, \eta) = \begin{cases} \frac{-1}{2} \|\eta\|_2^2 - {}^t\eta y & \text{if } \|{}^t X \eta\|_\infty \leq \lambda \\ -\infty & \text{otherwise} \end{cases}$$

So the dual of (LASSO) is :

$$\begin{aligned} &\text{maximize w.r.t. } \eta \quad \frac{-1}{2} \|\eta\|_2^2 - {}^t\eta y \\ &\text{subject to} \quad \|{}^t X \eta\|_\infty \leq \lambda \end{aligned}$$

And renaming η for v , multiplying by -1 to transform the maximization problem into a minimization problem, and reformulating $\|{}^tX\eta\|_\infty \leq \lambda$ as ${}^tX\eta \preceq b$ and $-{}^tX\eta \preceq b$ with $b = \begin{pmatrix} 1 \\ \dots \\ 1 \end{pmatrix} \lambda$:

$$\begin{array}{ll} \text{minimize w.r.t. } v & \frac{1}{2} {}^t v v + {}^t y v \\ \text{subject to} & {}^t X v \preceq b \\ & -{}^t X v \preceq b \end{array}$$

And therefore :

$$\begin{array}{ll} \text{minimize w.r.t. } v & {}^t v Q v + {}^t p v \\ \text{subject to} & A v \preceq b \end{array} \quad \text{with} \quad \begin{cases} Q = \frac{1}{2} I \in \mathbb{R}^{n \times n} \\ p = y \in \mathbb{R}^n \\ A = \begin{pmatrix} {}^t X \\ -{}^t X \end{pmatrix} \in \mathbb{R}^{2d \times n} \\ b = \begin{pmatrix} 1 \\ \dots \\ 1 \end{pmatrix} \lambda \in \mathbb{R}^{2d} \end{cases} \quad (QP)$$

Question 2 : Implement the barrier method to solve QP

Let $m = 2d$.

The barrier method reformulates the inequality constraints using logarithmic barrier functions :

$$\text{minimize w.r.t. } v \quad {}^t [{}^t v Q v + {}^t p v] - \sum_{i=1}^m \log(b_i - {}^t a_i v) \quad (CP)$$

where $b = \begin{pmatrix} b_1 \\ \dots \\ b_m \end{pmatrix} \in \mathbb{R}^d$ and $A = \begin{pmatrix} {}^t a_1 \\ \dots \\ {}^t a_m \end{pmatrix} \in \mathbb{R}^{d \times n}$

Step 1 : write `centering_step (Q, p, A, b, t, v0, eps)` to solve (CP).

As there is no equality constraint, the easiest option is to use the regular Method.

```
In [ ]: def backtracking_line_search(f, grad_f_x, x, delta_x, alpha, beta, t0=1):
    assert 0 < alpha and alpha < .5
    assert 0 < beta and beta < 1
    fx = f(x)
    assert not np.isnan(fx)

    t = [t0]
    slope = alpha * np.dot(grad_f_x, delta_x)

    while True:
        y = f(x + t[-1] * delta_x)
        if not np.isnan(y):
            if f(x + t[-1] * delta_x) <= fx + t[-1] * slope:
                return np.array(t)
            t.append( beta * t[-1] )

    ### Unit test ###
    def unit():
        alpha, beta = 0.3, 0.8
        plt.title(rf"Backtracking lines search unit test with $\alpha={alpha}$ and $\beta={beta}$")
        T = np.linspace(-2, 2, 1000)
```

```

# f
f = lambda x : np.square(x) - np.log(1 - x)
plt.plot(T, f(T), label="f")

# x
x = -1.5
plt.scatter([x],[f(x)], label="$x$")

# Tangents
grad_f_x = 2*x + 1/(1 - x)
tg = lambda t : f(x) + grad_f_x * (t - x)
tg_alpha = lambda t : f(x) + alpha * grad_f_x * (t - x)
plt.plot(T, tg(T), label="Tangent")
plt.plot(T, tg_alpha(T), label=r"$\alpha$ Tangent")

# t0
delta_x = 1
t0 = backtracking_line_search(
    f=f,
    grad_f_x=grad_f_x,
    x=x,
    delta_x=delta_x,
    alpha=alpha,
    beta=beta,
    t0=3
)
plt.scatter(x + delta_x * t0, f(x + delta_x * t0), label="$t_0$")

# Plot
plt.legend()
plt.ylim([-0.1, 4.1])
plt.show()

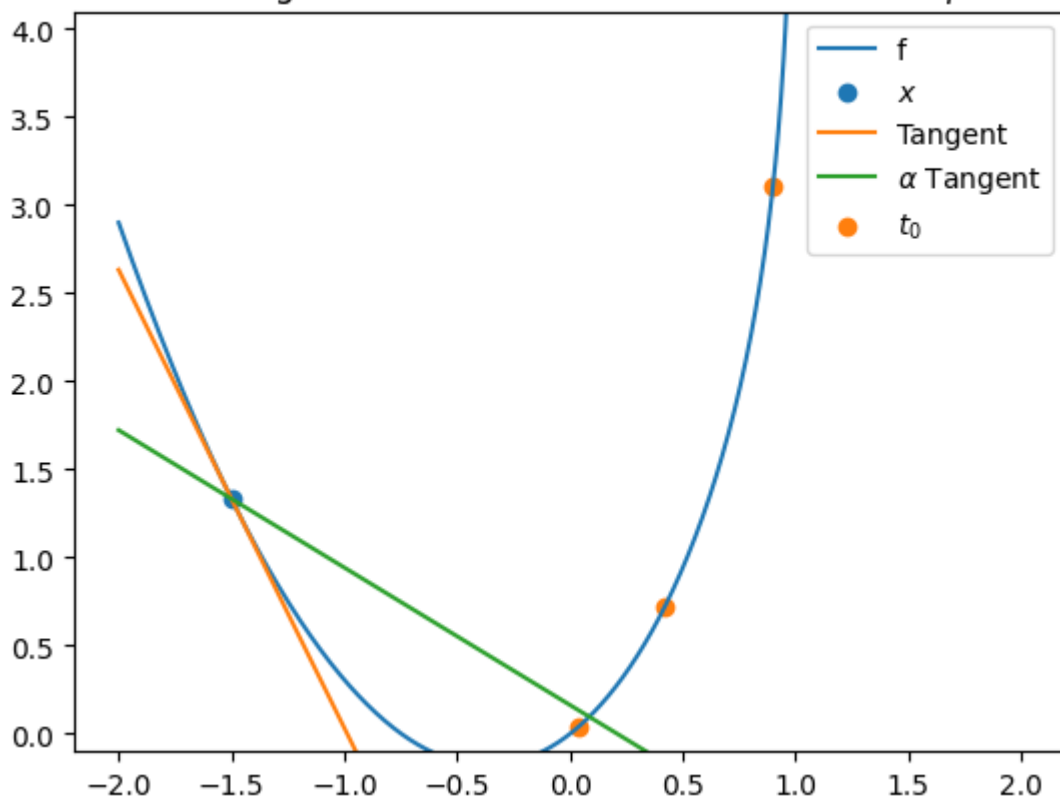
unit()

```

C:\Users\Romain\AppData\Local\Temp\ipykernel_4840\2032140524.py:24: RuntimeWarning: invalid value encountered in log

```
f = lambda x : np.square(x) - np.log(1 - x)
```

Backtracking lines search unit test with $\alpha = 0.3$ and $\beta = 0.8$



As this function will be extensively used, here is a optimized version which does not create a list at each

call, leading to a general increase in performances :

```
In [ ]: def faster_backtracking_line_search(f, grad_f_x, x, delta_x, alpha, beta, t0=1):
    assert 0 < alpha and alpha < .5
    assert 0 < beta and beta < 1
    fx = f(x)
    assert not np.isnan(fx)

    t = t0
    slope = alpha * np.dot(grad_f_x, delta_x)

    while True:
        y = f(x + t * delta_x)
        if not np.isnan(y):
            if f(x + t * delta_x) <= fx + t * slope:
                return t
            t = beta * t

def test_backtracking_function(backtracking_function, N = 10000):
    alpha, beta = 0.3, 0.8
    f = lambda x : np.square(x) - np.log(1 - x)
    x = -1.5
    grad_f_x = 2*x + 1/(1 - x)
    delta_x = 1
    t0 = 3
    for _ in range(N):
        backtracking_function(
            f=f,
            grad_f_x=grad_f_x,
            x=x,
            delta_x=delta_x,
            alpha=alpha,
            beta=beta,
        )
```

```
In [ ]: %%time
test_backtracking_function(backtracking_function=backtracking_line_search)
```

CPU times: total: 93.8 ms
Wall time: 94.4 ms

```
In [ ]: %%time
test_backtracking_function(backtracking_function=faster_backtracking_line_search)
```

CPU times: total: 93.8 ms
Wall time: 89.7 ms

Let f be :

$$f : v \mapsto \underbrace{t[vQv + {}^tpv]}_{f_0(v)} - \sum_{i=1}^d \underbrace{\log(b_i - {}^ta_iv)}_{-\phi(v)} = tf_0(v) + \sum_{i=1}^d \phi_i(v)$$

```
In [ ]: def build_f0(Q, p):
    return lambda u : np.dot(Q @ u + p, u)

def build_f(Q, p, A, b, t):
    f0 = build_f0(Q=Q, p=p)
    return lambda u : t * f0(u) - np.sum( np.log( b - A @ u ) )
```

Then its gradient and hessian are :

$$\nabla f = t \nabla f_0 + \sum_{i=1}^d \nabla \phi_i ; \quad \nabla^2 f = t \nabla^2 f_0 + \sum_{i=1}^d \nabla^2 \phi_i$$

f_0 is just a quadratic function, so its gradient and hessian are straight forward to compute (Q is assumed symmetric positive semi definite):

$$\nabla f_0(v) = 2Qv + p ; \quad \nabla^2 f_0(v) = 2Q$$

```
In [ ]: def f0_derivatives(Q, p, v):
        return 2 * Q @ v + p, 2 * Q
```

ϕ_i is the log barrier function with affine constraints :

$$\phi_i(v) = -\log(b_i - {}^t a_i v)$$

$$\nabla \phi_i(v) = \frac{1}{b_i - {}^t a_i v} a_i$$

$$\nabla^2 \phi_i(v) = \frac{1}{(b_i - {}^t a_i v)^2} a_i {}^t a_i = \nabla \phi_i(v) {}^t \nabla \phi_i(v)$$

To compute those efficiently, it is better to limit the amount of for loop and to use tensor operations.

First, $\text{affine} = \begin{pmatrix} b_1 - {}^t a_1 v \\ \dots \\ b_d - {}^t a_d v \end{pmatrix} = b - Av$ given that $A = \begin{pmatrix} {}^t a_1 \\ \dots \\ {}^t a_d \end{pmatrix}$.

Using numpy's elementwise inversion, it is easy to define $1/\text{affine} = \begin{pmatrix} \frac{1}{b_1 - {}^t a_1 v} \\ \dots \\ \frac{1}{b_d - {}^t a_d v} \end{pmatrix}$

Then $\text{grad} = (\nabla \phi_1(v) \dots \nabla \phi_d(v)) = \begin{pmatrix} a_1 \frac{1}{b_1 - {}^t a_1 v} & \dots & a_n \frac{1}{b_d - {}^t a_d v} \end{pmatrix}$, this is a columnwise multiplication (or multiplication by a diagonal matrix), so $\nabla \phi_i(v) = {}^t A \text{diag} \begin{pmatrix} \frac{1}{b_1 - {}^t a_1 v} \\ \dots \\ \frac{1}{b_d - {}^t a_d v} \end{pmatrix}$

And finally $\nabla^2 \phi_i(v) = \nabla \phi_i(v) {}^t \nabla \phi_i(v)$, this is a batchify outer product, hence $\text{hess} = \text{einsum}(\text{"nd,md->mnd"}, \text{grad}, \text{grad})$ (m and n has no importance as the hessian is symmetric).

And given the sum in the final expression, there is left to sum over the d -dimensional axis for both gradients and Hessians.

```
In [ ]: def phi_derivatives(A, b, v):
        affine = b - A @ v
        grad = A.T @ np.diag(1/affine)
        hess = np.einsum("nd,md->mnd", grad, grad)
        return np.sum(grad, axis=-1), np.sum(hess, axis=-1)

### Unit test ###
def unit():
    d=2
    n=3
    A = np.arange(n*d).reshape(n,d).T
    b = np.arange(n*d + 1, n*d + 1 + d)
    v = np.linspace(0, 1, n)

    return A, b, v, *phi_derivatives(A=A, b=b, v=v)
```

```
# This result is indeed the one I found by hand calculation
unit()
```

```
Out[ ]: (array([[0, 2, 4],
               [1, 3, 5]]),
        array([7, 8]),
        array([0. , 0.5, 1. ]),
        array([0.66666667, 3.          , 5.33333333]),
        array([[ 0.44444444,  1.33333333,  2.22222222],
               [ 1.33333333,  5.          ,  8.66666667],
               [ 2.22222222,  8.66666667, 15.11111111]]))
```

This leads to the following implementation of the newton method :

```
In [ ]: def centering_step(Q, p, A, b, t, v0, eps=1e-7, alpha=0.1, beta=0.9):
        # v0 must be in the domain of the Log,
        # ie a feasible point in the original problem
        assert np.all(b - A @ v0 > 0)
        v = [v0]
        f = build_f(Q=Q, p=p, A=A, b=b, t=t)

        while True:
            if len(v) % 50 == 0:
                print(f"Iteration {len(v)} : 0.5*squared_decrement={squared_decrement}")
            phi_grad, phi_hessian = phi_derivatives(A=A, b=b, v=v[-1])
            f0_grad, f0_hessian = f0_derivatives(Q=Q, p=p, v=v[-1])
            grad = t * f0_grad + phi_grad
            hessian = t * f0_hessian + phi_hessian

            delta_v = np.linalg.solve(hessian, -grad)
            squared_decrement = np.dot(grad, -delta_v)

            if 0.5*squared_decrement <= eps: return v
            s = faster_backtracking_line_search(
                f=f, x=v[-1],
                grad_f_x=grad,
                delta_x=delta_v,
                alpha=alpha, beta=beta
            )

            v.append( v[-1] + s * delta_v )
```

```
In [ ]: ### Unit Test ###
def unit():
    N = 100
    Q = np.diag((1,2))
    p = -2*np.arange(1, 2)
    A = np.ones(shape=(1,2))
    b = np.array([1])
    t = 0.05

    f = build_f(Q=Q, p=p, A=A, b=b, t=t)
    X, Y = np.linspace(-6, 6, N), np.linspace(-6, 6, N)
    Z = np.zeros((N, N))

    for j, y in enumerate(X):
        for i, x in enumerate(Y):
            v = np.array((x,y))
            Z[j,i] = f(v)

    fig = plt.figure()
    ax = plt.axes(projection='3d')
```

```

X, Y = np.meshgrid(X, Y)
ax.contour3D(X, Y, Z, 30, cmap='viridis')

V = centering_step(
    Q=Q, p=p,
    A=A, b=b,
    t=t,
    v0=-5*np.ones(shape=(2,))
)

for v in V:
    ax.scatter([v[0]], [v[1]], [f(v)], s=100, c='red')

print(f"v* = {V[-1]}, f(v*) = {f(V[-1]):0.5f}")
# ax.view_init(90, 0)
fig.show()

unit()

```

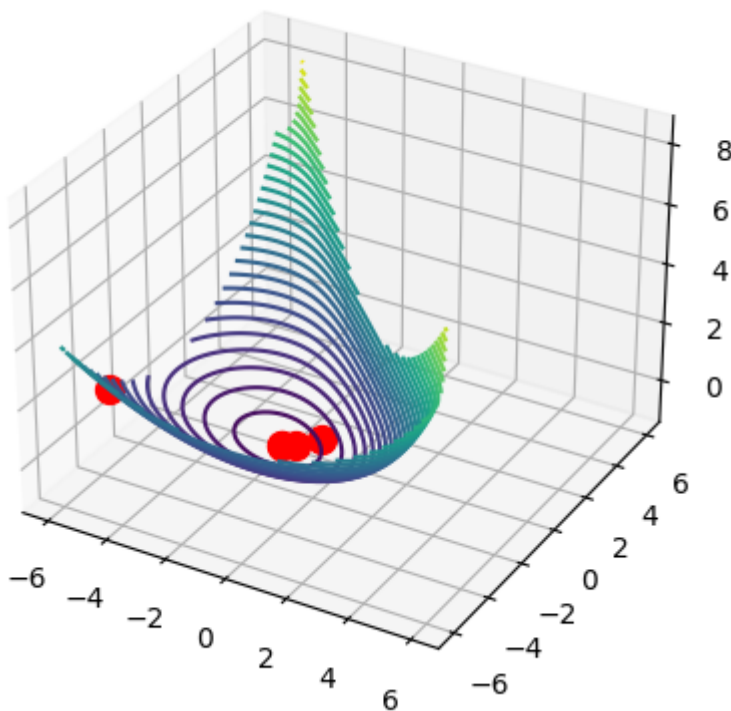
C:\Users\Romain\AppData\Local\Temp\ipykernel_4840\3029965375.py:6: RuntimeWarning: invalid value encountered in log

```
return lambda u : t * f0(u) - np.sum( np.log( b - A @ u ) )
```

C:\Users\Romain\AppData\Local\Temp\ipykernel_4840\1343674766.py:37: UserWarning: Matplotlib is currently using module://matplotlib_inline.backend_inline, which is a non-GUI backend, so cannot show the figure.

```
fig.show()
```

```
v* = [-1.7536703 -0.87683515], f(v*) = -0.79567
```



Step 2 : write `barr_method (Q, p, A, b, v0, eps)` to solve (QP).

```

In [ ]: def barr_method (Q, p, A, b, v0, mu, t0=1, eps=1e-6, alpha=0.3, beta=0.8):
    v, t = [v0], t0
    m = b.shape[0]

    for _ in tqdm(range( int(np.log(m/(eps*t0))/np.log(mu)) + 1), desc=f"Barrier method for m
    v += centering_step(
        Q=Q, p=p,
        A=A, b=b,
        t=t, v0=v[-1],
        eps=eps,
        alpha=alpha, beta=beta
    )

```

```

t = t*mu

return v

```

Question 3 : Test your function on randomly generated matrices X and observations y

```

In [ ]: def Generate_Data (n, d, noise_std=0.5):
        X = 2*np.random.rand(n, d) - 1
        w = 2*np.random.rand(d) - 1
        y = X @ w + np.random.normal(scale=noise_std, size=(n,))
        return X, y, w

```

`format_problem` is a function formatting (LASSO) dual problem as a Quadratic Programm. As a reminder :

$$\begin{array}{ll}
 \text{minimize w.r.t. } v & {}^t v Q v + {}^t p v \\
 \text{subject to} & A v \preceq b
 \end{array}
 \quad \text{with} \quad
 \begin{cases}
 Q = \frac{1}{2} I \in \mathbb{R}^{n \times n} \\
 p = y \in \mathbb{R}^n \\
 A = \begin{pmatrix} {}^t X \\ -{}^t X \end{pmatrix} \in \mathbb{R}^{2d \times n} \\
 b = \begin{pmatrix} 1 \\ \dots \\ 1 \end{pmatrix} \lambda \in \mathbb{R}^{2d}
 \end{cases}
 \quad (QP)$$

Given that $\lambda > 0$, $v_0 = 0$ is always a feasible point.

```

In [ ]: def format_problem(X, y, lambd):
        # Returns Q, p, A, b, v0 feasible (= 0)
        n, d = X.shape
        return 0.5 * np.eye(n), y, np.concatenate((X.T, -X.T), axis=0), np.ones(shape=(2*d,)) *

```

```

In [ ]: def test_method(n, d, lambd, list_mu):
        ax = plt.gca()
        for mu in list_mu:
            X, y, _ = Generate_Data (n=n, d=d)
            Q, p, A, b, v0 = format_problem(X=X, y=y, lambd=lambd)
            v = barr_method (Q=Q, p=p, A=A, b=b, v0=v0, mu=mu, eps=1e-6)

            f0 = build_f0(Q=Q, p=p)
            v_star = v[-1]
            d_star = f0(v_star)

            ax.semilogy(np.arange(len(v)), [f0(u) - d_star for u in v], label=fr"$\mu={mu}$")

        ax.set_title(rf"$f_0(u) - d^*$ overtime for $\lambda = {lambd}$.")
        ax.grid(True, which="both")
        ax.set_xlabel(r"$N_{\text{steps}}$")
        plt.legend()
        plt.show()

```

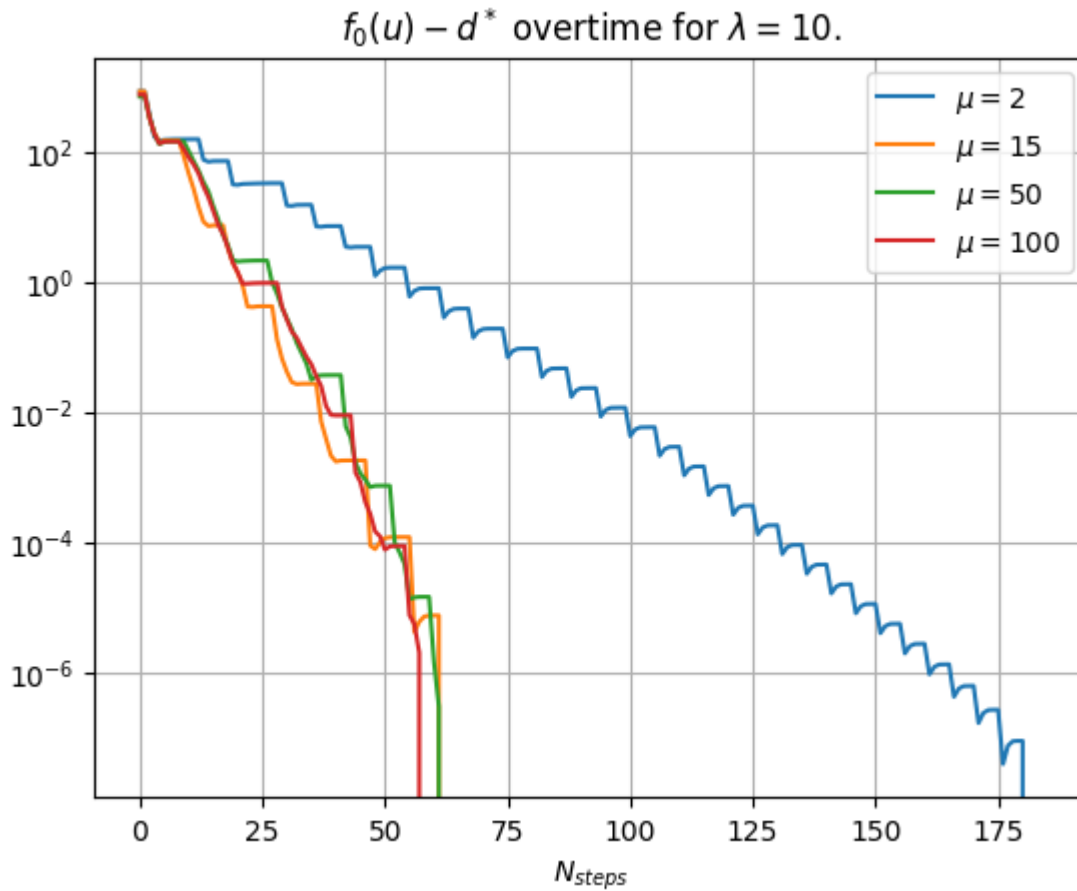
```

test_method(n=100, d=1000, lambd=10, list_mu=[2, 15, 50, 100])

```



```
Barrier method for mu = 2: 0%| 0/31 [00:00<?, ?it/s]C:\Users\Romain\AppData\Local\Temp\ipykernel_4840\3029965375.py:6: RuntimeWarning: invalid value encountered in log
    return lambda u : t * f0(u) - np.sum( np.log( b - A @ u ) )
Barrier method for mu = 2: 100%| 31/31 [00:18<00:00, 1.71it/s]
Barrier method for mu = 15: 100%| 8/8 [00:06<00:00, 1.23it/s]
Barrier method for mu = 50: 100%| 6/6 [00:06<00:00, 1.09s/it]
Barrier method for mu = 100: 100%| 5/5 [00:06<00:00, 1.25s/it]
```



This matches the theoretical results. Now it is time to retrieve w^* from the solution of the dual.

The original problem (LASSO) is strictly feasible (there is no inequality constraints and the objective function is defined on \mathbb{R}^d) and the problem is convex, hence **Slater's conditions** are gathered and **strong duality holds**.

It means that the KKT's conditions are respected. Therefore, the lagrangien at the optimum $\mathcal{L}(w^*, z^*, v^*)$ is minimal.

As a reminder :

$$\mathcal{L}(w, z, v) = \underbrace{\frac{1}{2} \|z\|_2^2 - {}^t v z}_{f(z)} + \underbrace{\lambda \|w\|_1 + {}^t v X w - {}^t v y}_{g(w)}$$

First, \mathcal{L} is differentiable with respect to z , and $\nabla_z \mathcal{L}(z^*, w^*, v^*) = 0 \Leftrightarrow z^* = v^*$. z is by construction equal to $Xw - y$, hence $Xw^* = y + v^*$.

This equation is equivalent to ${}^t X X w^* = {}^t X (y + v^*)$. But $X \in \mathbb{R}^{n \times d}$, hence X is full rank (and the equation can be solve) only if $n \geq d$. But LASSO is designed to tackle this problem in a high dimensional setting, when $d \gg n$, so this alone is not enough to solve in this setting.

\mathcal{L} is not differentiable with respect to w due to the absolute value function. However, the l_1 -norm admits subderivatives. Hence \mathcal{L} has subderivatives and the relaxed KKT's condition states that at least one of those subderivative is 0 at the optimum.

A subderivative of $\|\cdot\|_1$ is a function that maps any x in \mathbb{R}^d to a vector $y \in \mathbb{R}^d$ such that for

$$1 \leq i \leq n, y_i = \begin{cases} 1 & \text{if } x_i > 0 \\ -1 & \text{if } x_i < 0 \\ \alpha & \text{if } x_i = 0 \end{cases} \text{ for a given } \alpha \in [-1, 1]. \text{ This function is note } \nabla^{(\alpha)} \|\cdot\|_1.$$

Let use the KKT conditions and note $\nabla_w \mathcal{L}$ a subderivative of \mathcal{L} with respect to w , defined with an $\alpha \in [-1, 1]$ such that $\nabla_w \mathcal{L}(z^*, w^*, v^*) = 0$.

This means that $\lambda \nabla^{(\alpha)} \|w^*\|_1 + {}^t X v^* = 0$, which is equivalent to :

$$\nabla^{(\alpha)} \|w^*\|_1 = -\frac{1}{\lambda} {}^t X v^*$$

Then if $w_i^* \neq 0$, then $|\frac{1}{\lambda} {}^t X v^*|_i = 1$, the contraposition being $|\frac{1}{\lambda} {}^t X v^*|_i < 1 \Rightarrow w_i^* = 0$ (Indeed, the dual's inequality constraint imposes $\|\frac{1}{\lambda} {}^t X v^*\|_1 \leq 1$).

This criteria helps to find the 0 of w , and if it sets enough 0, then the original system might become solvable.

Here is an implementation of this method to find w . Note that the checked condition is relaxed with a parameter $\epsilon = 10^{-6}$, because the barrier method introduces a slight offset. It becomes $|\frac{1}{\lambda} {}^t X v^*|_i < 1 - \epsilon \Rightarrow w_i^* = 0$.

```
In [ ]: def retrieve_w_from_v(v, X, y, lambd, eps= 1e-6):
    n, d = X.shape
    w = np.zeros((d,))
    crit = X.T @ v
    non_zeros = np.abs(crit) >= lambd - eps
    print(f"Found {d - np.sum(non_zeros)} zeros in w for n = {X.shape[0]} and d = {w.shape[0]}")

    shorten_X = X[:, non_zeros]
    w[non_zeros] = np.linalg.solve(shorten_X.T @ shorten_X, shorten_X.T @ (y + v))
    return w

def unit(n, d, lambd, list_mu):
    for mu in list_mu:
        X, y, _ = Generate_Data (n=n, d=d)
        Q, p, A, b, v0 = format_problem(X=X, y=y, lambd=lambd)
        v = barr_method (Q=Q, p=p, A=A, b=b, v0=v0, mu=mu, eps=1e-6)

        f0 = build_f0(Q=Q, p=p)
        v_star = v[-1]
        d_star = f0(v_star)

        w_star = retrieve_w_from_v(v=v_star, X=X, y=y, lambd=lambd)
        p_star = np.dot(X @ w_star - y, X @ w_star - y) + lambd * np.sum(np.abs(w_star))

        print(f"d_star = {-d_star:0.1f}, p_star = {p_star:0.1f}, dual_gap = {(p_star + d_star) / 2 : 0.1f}")

unit(n=100, d=1000, lambd=10, list_mu=[2, 15, 50, 100, 500, 1000])
```

```
Barrier method for mu = 2: 0%|██████████| 0/31 [00:00<?, ?it/s]C:\Users\Romain\AppData\Local\Temp\ipykernel_4840\3029965375.py:6: RuntimeWarning: invalid value encountered in log
    return lambda u : t * f0(u) - np.sum( np.log( b - A @ u ) )
Barrier method for mu = 2: 100%|██████████| 31/31 [00:17<00:00, 1.81it/s]
Found 907 zeros in w for n = 100 and d = 1000.
d_star = 859.5, p_star = 909.0, dual_gap = 49.54, relative dual gap = 5.4%
Barrier method for mu = 15: 100%|██████████| 8/8 [00:07<00:00, 1.09it/s]
Found 908 zeros in w for n = 100 and d = 1000.
d_star = 810.6, p_star = 854.5, dual_gap = 43.84, relative dual gap = 5.1%
Barrier method for mu = 50: 100%|██████████| 6/6 [00:06<00:00, 1.05s/it]
```

Found 912 zeros in w for $n = 100$ and $d = 1000$.

$d_{\text{star}} = 827.4$, $p_{\text{star}} = 875.9$, $\text{dual_gap} = 48.55$, relative dual gap = 5.5%

Barrier method for $\mu = 100$: 100%|██████████| 5/5 [00:06<00:00, 1.37s/it]

Found 908 zeros in w for $n = 100$ and $d = 1000$.

$d_{\text{star}} = 990.7$, $p_{\text{star}} = 1044.5$, $\text{dual_gap} = 53.85$, relative dual gap = 5.2%

Barrier method for $\mu = 500$: 100%|██████████| 4/4 [00:08<00:00, 2.04s/it]

Found 907 zeros in w for $n = 100$ and $d = 1000$.

$d_{\text{star}} = 798.0$, $p_{\text{star}} = 846.2$, $\text{dual_gap} = 48.21$, relative dual gap = 5.7%

Barrier method for $\mu = 1000$: 100%|██████████| 4/4 [00:08<00:00, 2.05s/it]

Found 909 zeros in w for $n = 100$ and $d = 1000$.

$d_{\text{star}} = 829.4$, $p_{\text{star}} = 875.9$, $\text{dual_gap} = 46.53$, relative dual gap = 5.3%

It looks like a higher μ does not change the performances of the found w . Given that it takes fewer iterations to converge, it looks reasonable to take a high value for μ .